

# Inferring Harmony from Free Polyphony

Judah Daniels

April 5, 2023

DOS: Prof. Larry Paulson

Crsid: jasd6

College: Clare College

## 1 Abstract

A piece of music can be described using a sequence of chords, representing a higher level harmonic structure of a piece. There is a small, finite set of chord types, but each chord can be realised on the musical surface in a practically infinite number of ways. Given a score, we wish to infer the underlying chord types.

The paper *Modeling and Inferring Proto-voice Structure in Free Polyphony* describes a generative model that encodes the recursive and hierarchical dependencies between notes, giving rise to a grammar-like hierarchical system [2]. This proto-voice model can be used to reduce a piece into a hierarchical structure which encodes an understanding of the tonal/harmonic relations of a piece.

Christoph Finkensiep suggests in his paper that the proto-voice model may be an effective way to infer higher level latent entities, such as harmonies or voice leading schemata. Thus in this project I will ask the question: is this parsing model an effective way to annotate harmonies? By ‘effective’ we are referring to two things:

- Accuracy: can the model successfully emulate how experts annotate harmonic progressions in musical passages?
- Practicality: can the model be used to do this within a reasonable time frame?

While the original model could in theory be used to generate harmonic annotations, its exhaustive search strategy would be prohibitively time-consuming in practice for any but the shortest musical extracts; one half measure can have over 100,000 valid derivations [2]. My approach will be to explore the use of heuristic search algorithms to solve this problem.

## 2 Substance and Structure

### 2.1 Core: Search

The core of this project is essentially a search problem characterised as follows:

- The state space  $S$  is the set of all possible partial reductions of a piece along with each reduction step that has been done so far.
- We have an initial state  $s_o \in S$ , which is the empty reduction, corresponding to the unreduced surface of the piece. The score is represented as a sequence of slices grouping notes that sound simultaneously. We are also given the segmentation of the original chord labels that we wish to retrieve.
- We have a set of actions,  $A$  modelled by a function  $action : A \times S \rightarrow S$ . These actions correspond to a single reduction step.
  - The reduction steps are the inverses of the operations defined by the generative proto-voice model.
- Finally we have a goal test,  $goal : S \rightarrow \{true, false\}$  which is true iff the partial reduction  $s$  has exactly one slice per segment of the input.

- This means the partial reduction  $s$  contains a sequence of slices which start and end positions corresponding to the segmentation of the piece.
- At the first stage, this will be implemented using a random graph search algorithm, picking each action randomly, according to precomputed distributions.

## 2.2 Core: Evaluation

The second core task is to create an evaluation module that iterates over the test dataset, and evaluates the partial reduction computed by the search algorithm above. This will be done by comparing the outputs to ground truth annotations from the Annotated Beethoven Corpus.

In order to do this I will make use of the statistical harmony model from Finkensiep’s thesis, *The Structure of Free Polyphony* [1]. This model provides a way of mapping between the slices that the algorithm generates and the chords in the ground truth. This can be used to empirically measure how closely the slices match the expert annotations.

## 2.3 Extension

Once the base search implementation and evaluation module have been completed, the search problem will be tackled by heuristic search methods, with different heuristics to be trialled and evaluated against each other. The heuristics will make use of the chord profiles from Finkensiep’s statistical harmony model discussed above. These profiles relate note choices to the underlying harmony. Hence the heuristics may include:

- How the chord types relate to the pitches used.
- How the chord types relate which notes are used as ornamentation, and the degree of ornamentation.
- Contextual information about neighboring slices

## 2.4 Overview

The main work packages are as follows:

**Preliminary Reading** – Familiarise myself with the proto-voice model, and read up on similar models and their implementations. Study heuristic search algorithms.

**Dataset Preparation** – Pre-process the Annotated Beethoven Corpus into a suitable representation for my algorithm.

**Basic Search** – Implement a basic random search algorithm that takes in surface and segmentations, and outputting the sequence of slices matching the segmentations.

**Evaluation Module** – Implement an evaluation module to evaluate the output from the search algorithm.

**End-to-end pipeline** – Implement a full pipeline from the data to the evaluation that can be used to compare different reductions.

**Heuristic Design** – Extension – Trial different heuristics and evaluate their performance against each other.

**Dissertation** – I intend to work on the dissertation throughout the duration of the project. I will then focus on completing and polishing the project upon completion.

# 3 Starting Point

The following describes existing code and languages that will be used for this project:

**Haskell** – I will be using Haskell for this project as it is used in the proto-voice implementation. It must be noted that my experience with Haskell is limited, as I was first introduced to it via an internship this summer (July to August 2022).

**Python** – Python will be used for data handling. I have experience coding in Python.

**Prior Research** - Over the summer I have been reading the literature on computational models of music, as well as various parsing algorithms such as semi-ring parsing [3], and the CYK algorithm, which is used in the implementation of the proto-voice model.

**Protovoices-Haskell** - The paper *Modeling and Inferring Proto-Voice Structure in Free Polyphony* [2] includes an implementation of the proto-voice model in Haskell. A fork of this repository will form the basis of my project. This repository includes a parsing module which will be used to perform the actions in the search space of partial reductions. There is a module that can exhaustively enumerate reductions of a piece, but this is infeasible in practice due to the blowup of the derivation forest.

**MS3** - This is a library for parsing MuseScore Files and manipulating labels [4], which I will use as part of the data processing pipeline.

**ABC** - The *Annotated Beethoven Corpus* [5] contains analyses of all Beethoven string quartets composed between 1800 and 1826, encoded in a human and machine readable format. This will be used as a dataset for this project.

## 4 Success Criteria

This project will be deemed a success if I complete the following tasks:

- Develop a baseline search algorithm that uses the proto-voice model to output a partial reduction of a piece of music up to the chord labels.
- Create an evaluation module that can take the output of the search algorithm and quantitatively evaluate its accuracy against the ground truth annotations by providing a score based on a statistical harmony model.
- Extension: Develop one or more search algorithms that use additional heuristics to inform the search, and compare the accuracy with the baseline algorithm.

## 5 Timetable

Time frame	Work	Evidence
<b>Michaelmas (Oct 4 to Dec 2)</b>		
Oct 14 to Oct 24	<i>Oct 14</i> : Final proposal deadline. Preparation work: familiarise myself with the dataset and the proto-voice model implementation. Work on manipulating reductions using the proto-voice parser provided by the paper.	None
Oct 24 to Nov 7	Dataset preparation and handling.	Plot useful metrics about the dataset using Haskell
Nov 7 to Nov 21	Random Search implementation	None
Nov 21 to Dec 5	Evaluation Module. Continue with search implementation.	Evaluate a manually created derivation and plot results
<b>Vacation (Dec 3 to Jan 16)</b>		
Dec 5 to Dec 11	Evaluate performance of random search. Begin to work on extensions	Plot results
Dec 10 to Dec 21	Trial different heuristics. Implement an end-to-end pipeline from input to evaluation.	None
Dec 21 to Dec 27	None	None
Dec 27 to Jan 10	Continue trialing and evaluating heuristics	<i>Fulfill success criterion: At least one heuristic technique gives better performance than random search.</i>
<b>Lent (Jan 17 to Mar 17)</b>		
Jan 4 to Jan 20	Buffer Period to help keep on track	None
Jan 20 to Feb 3	<i>Feb 3</i> : Progress Report Deadline. Write progress report and prepare presentation. Write draft <i>Evaluation</i> chapter	Progress Report (approx. 1 page)
Feb 3 to Feb 17	Prepare presentation.	<i>Feb 8 – 15</i> : Progress Report presentation
Feb 17 to Mar 3	<i>Feb 17</i> : How to write a Dissertation briefing. Write draft Introduction and Preparation chapters. Incorporate feedback on Evaluation chapter.	Send draft Introduction and Preparation chapter to supervisor
Mar 3 to Mar 17	Write draft Implementation chapters. Incorporate feedback on Introduction and Preparation chapters.	Send draft Implementation chapters to Supervisor
<b>Vacation (Mar 18 to Apr 24)</b>		
Mar 18 to Mar 31	Complete draft dissertation.	Send draft dissertation to supervisor
Mar 31 to April 15	Give time for supervisor to read dissertation	None
April 15 to April 25	Incorporate feedback, make final revisions and checks of the whole dissertation	Submit dissertation and source code
<b>Easter (Apr 25 to Jun 16)</b>		
April 25 to May 12	None	None

## 6 Resources

I plan to use my own laptop for development: MacBook Pro 16-inch, M1 Max, 32GB Ram, 1TB SSD, 24-core GPU.

All code will be stored on a GitHub repository, which will guarantee protection from data loss. I will easily be able to switch to using university provided computers upon hardware/software failure.

The project will be built upon work that has been done in the DCML (Digital cognitive musicology lab) based in EPFL. The files are in their Github repository, and I have been granted permission to access their in-house datasets of score annotations, as well as software packages which are used to handle the data.

## 7 Supervisor Information

Peter Harrison, head of Centre for Music and Science at Cambridge, has agreed to supervise me for this. We have agreed on a timetable for supervisions for this year. I am also working with Christoph Finkensiep, a PHD student at the DCML, and originator of the proto-voice model. Professor Larry Paulson has agreed to be the representative university teaching officer.

## References

- [1] Christoph Finkensiep. *The Structure of Free Polyphony*. PhD thesis, EPFL, Lausanne, 2023.
- [2] Christoph Finkensiep and Martin Rohrmeier. Modeling and Inferring Proto-voice Structure in Free Polyphony. In *Proceedings of the 22nd ISMIR Conference*, Online, November 2021.
- [3] Joshua Goodman. Semiring Parsing. *Computational Linguistics*, 25(4):573–606, 1999.
- [4] Hentschel Johannes. Ms3 - Parsing MuseScore 3. <https://github.com/johentsch/ms3>, 2021.
- [5] Markus Neuwirth, Daniel Harasim, Fabian C. Moss, and Martin Rohrmeier. The Annotated Beethoven Corpus (ABC): A Dataset of Harmonic Analyses of All Beethoven String Quartets. *Frontiers in Digital Humanities*, 5:16, July 2018.